

Hybrid Search in PostgreSQL

Combining Vector and Full-Text for Real-World Applications

Shayon Sanyal · Principal PostgreSQL Specialist SA · Lead, Agentic AI for Databases

Postgres Summit US 2026 · New York City
September 30 · 10:30–11:20 EDT · Letterpress

Two questions. Two different misses.

Keyword search misses

"Where should I park my rainy-day / emergency fund?"

The best answer is **#2** for vector search. Keyword search ranks it **#29**, BM25 **#40**, and requiring every word finds nothing.

Vector search misses

"Employer rollover from 403b to 401k?"

The only answer is **#1** for BM25 and **#26** for vector search. Exact identifiers are what keyword search is for.

Words find exact terms. Vectors find paraphrases. Which one wins, and whether combining them helps, is a measurement, not a belief.



About me

Shayon Sanyal

Principal PostgreSQL Specialist Solutions Architect

Tech Lead — Agentic AI for Databases

I help teams build on PostgreSQL, from relational applications to retrieval and agent workflows.

Today: readable SQL, measured results, and a skill you can point at your own tables.

linkedin.com/in/shayonsanyal

How we'll measure

2,271 questions with known right answers.


FiQA-2018: 57,638 finance forum posts, 648 test questions. The main example.

SciFact, NFCorpus, SCIDOCS: chosen by one rule before measuring: BEIR sets under 30,000 documents where BM25 beat every dense retriever in 2021.

NDCG@10 scores a ranking higher when the known answers sit near the top.

Credit for a known answer at each rank

$1 / \log_2(\text{rank} + 1)$: an answer at rank 1 is worth more than three answers at rank 10



The stack: one PostgreSQL database

Piece	Version	Job
PostgreSQL	18.6	<code>tsvector</code> , GIN, <code>ts_rank_cd</code> , the fusion SQL, the scoring SQL
pgvector	0.8.6	<code>vector(1536)</code> , HNSW, iterative index scans, <code>halfvec</code> , <code>bit</code>
pg_textsearch	1.4.0	BM25 index and <code><@></code> operator (PostgreSQL license)
Cohere Embed v4 · Amazon Bedrock	1536 dims	Posts as <code>search_document</code> , questions as <code>search_query</code>
Cohere Rerank · Amazon Bedrock	3.5	Reorders a top 50 outside the database
bge-small-en-v1.5 · fastembed	384 dims	A small open-source model, on this laptop
VS Code + SQLTools		Numbered SQL files you can run yourself

Everything except the model calls is SQL. Every number today is from this laptop.

One row, three representations · 01_schema.sql

```
CREATE TABLE docs (  
  id          text PRIMARY KEY,  
  body        text NOT NULL,  
  tsv         tsvector GENERATED ALWAYS AS (to_tsvector('english', body)) STORED,  
  embedding   vector(1536), -- Cohere Embed v4, input_type = 'search_document'  
  embedding_local vector(384) -- bge-small-en-v1.5, open source, local  
);  
CREATE INDEX ON docs USING gin (tsv); -- 0.6 s  
CREATE INDEX ON docs USING hnsw (embedding vector_cosine_ops); -- 30.4 s  
CREATE INDEX ON docs USING bm25 (body) WITH (text_config = 'english'); -- 3.9 s
```

The lexemes are **generated** and cannot drift from the text. Embeddings are not: **one column per model**, re-embedded when the text changes.

Pitfall 1: every word must match · 03

```
SELECT websearch_to_tsquery('english', 'Where should I park my rainy-day / emergency fund?');  
-- 'park' & 'rainy-day' <=> 'raini' <=> 'day' & 'emerg' & 'fund'
```

`websearch_to_tsquery` and `plainto_tsquery` join words with **AND**. A natural question rarely has all its words in one answer.

405 of 648 questions match no post at all · NDCG@10 4.3

Match any word, rank the matches · 04

```
WITH question AS (  
    SELECT replace(plainto_tsquery('english', body)::text, ' & ', ' | ')::tsquery AS tsq  
    FROM active_query  
)  
SELECT d.id, ts_rank_cd(d.tsv, q.tsq) AS score  
    FROM docs d, question q  
WHERE d.tsv @@ q.tsq  
ORDER BY score DESC  
LIMIT 50;
```

Recall@50 doubles (6.7 → 14.6), but NDCG@10 **falls to 2.8**: with no IDF, “fund” and “day” outweigh “rainy-day”. The rainy-day question matches **10,460 posts**, and `ts_rank_cd` scores every one of them: **81 ms**.

ts_rank_cd is not BM25 · 05

```
SELECT id, -(body <@> to_bm25query(:'question', 'docs_body_bm25')) AS bm25
FROM docs
ORDER BY body <@> to_bm25query(:'question', 'docs_body_bm25')
LIMIT 50;
```

BM25 weighs rare terms up (IDF), saturates repeats, and normalizes for length. The index keeps the corpus statistics and returns the top 50 **without scoring every match**. 23.6 is exactly BEIR's published BM25 score for FiQA (0.236).

Keyword arm	NDCG@10	Top 50, rainy-day question
ts_rank_cd, any word	2.8	81 ms
BM25, pg_textsearch	23.6	0.27 ms

Semantic candidates · 06

```
SET hnsw.ef_search = 100;           -- default 40: LIMIT 50 would silently return 40
SELECT id, 1 - (embedding <=> (SELECT embedding FROM active_query)) AS cosine
FROM docs
WHERE embedding IS NOT NULL
ORDER BY embedding <=> (SELECT embedding FROM active_query)
LIMIT 50;
```

Keep the distance operator in `ORDER BY`, ascending, with a `LIMIT`. Embed questions as `search_query` and posts as `search_document`.

NDCG@10 53.9 · p50 3.5 ms

Pitfall 2: adding scores from different scales · 07a 07b

Rainy-day question, each list's top 50	min	max
ts_rank_cd	1.70	5.40
cosine similarity	0.478	0.598
Fusion	NDCG@10	
keyword_score + cosine		5.6
keyword list, then vector list, duplicates removed		2.9
Reciprocal Rank Fusion of the same two lists		33.8

Same inputs, 6× the score. Normalizing flags make ts_rank_cd bounded, not comparable.

Fuse ranks, not scores

$$\text{RRF}(d) = \sum 1 / (60 + \text{rank}_{\text{list}}(d))$$

"Transfer stock I own into my Roth IRA?" (FiQA 8512)	BM25 rank	Vector rank	RRF	Fused
The known answer	4	9	$1/64 + 1/69 = \mathbf{0.03012}$	#1
Another post	2	17	$1/62 + 1/77 = 0.02912$	#2
Another post	24	3	$1/84 + 1/63 = 0.02778$	#3
Vector's #1	—	1	$0 + 1/61 = 0.01639$	#9

Missing from a list contributes **0**. Agreement between lists wins; no score scales involved.

RRF is a full outer join · 08a

```
WITH keyword AS (  
    SELECT id, row_number() OVER (ORDER BY ts_rank_cd(tsv, q) DESC, id) AS rank  
    FROM docs, (SELECT replace(plainto_tsquery('english', :'question')::text,  
                                ' & ', ' | ')::tsquery AS q) question  
    WHERE tsv @@ q ORDER BY rank LIMIT 50  
)  
,  
semantic AS (  
    SELECT id, row_number() OVER (ORDER BY distance, id) AS rank  
    FROM (SELECT id, embedding <=> :'question_embedding' AS distance  
          FROM docs WHERE embedding IS NOT NULL  
          ORDER BY distance LIMIT 50) nearest  
)  
SELECT coalesce(k.id, v.id) AS id,  
       coalesce(1.0 / (60 + k.rank), 0) + coalesce(1.0 / (60 + v.rank), 0) AS rrf  
FROM keyword k FULL OUTER JOIN semantic v USING (id)  
ORDER BY rrf DESC  
LIMIT 10;
```

One statement, two indexes (GIN and HNSW). p50 **150 ms** with `ts_rank_cd` ; **8.2 ms** with BM25 as the keyword list (08b).

Blend scores the right way · 08c

```
keyword_norm AS (
    -- per question: min-max to [0, 1]; semantic_norm is the same
    SELECT id, (score - min(score) OVER ()) / (max(score) OVER () - min(score) OVER ()) AS s
    FROM keyword
)
SELECT id, w * coalesce(v.s, 0) + (1 - w) * coalesce(k.s, 0) AS blend
FROM keyword_norm k FULL OUTER JOIN semantic_norm v USING (id);
```

RRF ignores how **confident** each list is. A normalized blend keeps it, and is what “adding scores” should have been in pitfall 2.

Weight w on vector, tuned on dev questions only	FiQA	SciFact	NFCorpus	SCIDOCS
Cohere Embed v4 (frontier)	1.0	0.70	0.70	0.5, untuned
bge-small (local)	0.65	0.55	0.65	0.5, untuned

Knobs that change different things

Knob	Controls	Watch
Candidate depth (50)	How many posts each list contributes	A post outside both lists can't be fused or reranked: Recall@50
k (60)	How fast credit falls with rank	Smaller k favors each list's top few
Weights	Trust in each list	Equal-weight RRF lost to vector on all four datasets; tune a blend on dev questions instead
<code>hnsw.ef_search</code> (100)	Work per HNSW scan	Must be $\geq$ the candidate depth

Change one at a time, and re-run the test questions.

Live · NFCorpus · all local except the last column

Small model + BM25 beat the frontier model here

NFCorpus test question PLAIN-307, with 3 known answers · also active in VS Code

Vitamin D: Shedding some light on the new recommendations

A in 4 of 4 top 10s **B** in 3 of 4 top 10s **C** in 4 of 4 top 10s

BM25

SQL

65 NDCG@10

3 of 3 answers in the top 10
7.8 ms, live SQL

- 1 Safety considerations and potential interactions of vitamins: should vitamins be considered drugs?...
- 2 **B** VITAMIN D: A D-LIGHTFUL SOLUTION FOR HEALTH Throughout evolution sunlight produced vitamin D in the ski...
- 3 The aryl hydrocarbon receptor and its xenobiotic ligands: a fundamental

Vector · bge-small (local)

SQL

77 NDCG@10

2 of 3 answers in the top 10
17.9 ms, live SQL

- 1 **A** Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 2 **C** Evaluation, treatment, and prevention of vitamin D deficiency: an Endocrine Society clinical practice guideline....
- 3 Traditionally living populations in East

Tuned blend · bge-small

SQL

97 NDCG@10

Highest

3 of 3 answers in the top 10
19.2 ms, live SQL

- 1 **C** keyword 4 vector 2 Evaluation, treatment, and prevention of vitamin D deficiency: an Endocrine Society clinical practice guideline....
- 2 **A** keyword 8 vector 1 Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 3 keyword 14 vector 4

Vector · Embed v4

SQL

71 NDCG@10

3 of 3 answers in the top 10
25.9 ms, live SQL

- 1 Low Vitamin D Status: Definition, Prevalence, Consequences and Correction Vitamin D is obtained fro...
- 2 **A** Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 3 **B** VITAMIN D: A D-LIGHTFUL SOLUTION

localhost:8018 · pick NFCorpus · “Local hybrid beats the frontier model”

From text to vector

- **Cohere Embed v4** on Amazon Bedrock: 1536 dimensions, unit-length vectors, so cosine distance and inner product rank the same way.
- **Two input types.** Posts are embedded as `search_document`, questions as `search_query`. Swapping them lowers relevance without any error.
- **One model per column.** Vectors from different models or versions are not comparable, even at the same dimension. Store the model id.
- **A small local model too:** bge-small-en-v1.5, 384 dims, open source, in its own column. FiQA's 57,638 posts took about an hour on the laptop CPU; Embed v4 took 35 minutes behind a tokens-per-minute quota. Resume on NULLs either way.

Pitfall 3: a filter is not a pre-filter · 09

```
SELECT id FROM docs
WHERE tsv @@ plainto_tsquery('english', 'mortgage') -- 4.2% of posts
ORDER BY embedding <=> : 'question_embedding' LIMIT 10;
```

Setting	Plan	Rows returned
defaults (<code>iterative_scan = off</code> , <code>ef_search = 40</code>)	HNSW, then filter: Rows Removed by Filter: 39	1 of 10
<code>SET hnsw.iterative_scan = relaxed_order</code>	HNSW keeps walking until 10 pass	10 of 10
a rare term: <code>heloc</code> (0.2%)	GIN, then exact sort	10 of 10

pgvector 0.8 iterative scans are bounded by `hnsw.max_scan_tuples` (20,000).

Similar to the question AND contains a word · 11

```
SELECT h.*, left(d.body, 60)
  FROM hybrid_search(:'question', :'question_embedding',
                    match_count => 5, required_terms => 'mortgage') h
 JOIN docs d ON d.id = h.doc_id;
```

```
CREATE FUNCTION hybrid_search(...) ... LANGUAGE sql STABLE
SET hnsw.ef_search = 100
SET hnsw.iterative_scan = relaxed_order
SET plan_cache_mode = force_custom_plan -- keeps "required_terms IS NULL OR ..." indexable
```

The requirement goes into **both** candidate lists, before their LIMIT.

Hybrid vs vector, same embedding model · NDCG@10 change on 2,271 test questions

	FiQA	SciFact	NFCorpus	SCIDOCS
Small local model · bge-small, 384 dims, on this laptop				
Tuned blend + BM25	+1.2 ↑ +0.4...+2.1	+2.1 ↑ +0.0...+4.3	+2.1 ↑ +1.3...+3.0	+0.2 * -0.5...+0.8
Equal-weight RRF + BM25	-3.1 ↓ -4.7...-1.5	+1.7 -0.9...+4.3	+2.3 ↑ +1.0...+3.7	-0.3 -1.0...+0.5
Frontier model · Cohere Embed v4, 1536 dims				
Tuned blend + BM25	0.0 chose w = 1	+0.2 -1.2...+1.4	+0.8 ↑ +0.1...+1.5	-0.8 ↓ * -1.4...-0.2
Equal-weight RRF + BM25	-12.5 ↓ -14.5...-10.6	-2.9 ↓ -5.1...-0.6	-1.1 -2.2...+0.1	-1.2 ↓ -2.0...-0.5
Cohere Rerank 3.5, top 50	-4.1 ↓ -6.0...-2.3	-0.4 -2.8...+1.9	-1.9 ↓ -3.3...-0.4	-0.4 -1.1...+0.3

↑ / ↓: 95% paired bootstrap interval above / below zero · * SCIDOCS has no dev questions, so its blend weight is an untuned 0.5 · [results/summary.md](#)

When hybrid pays: a smaller model

NDCG@10	FiQA	SciFact	NFCorpus	SCIDOCS
bge-small alone · local, open source	38.0	72.0	33.8	19.6
bge-small + BM25, tuned blend · all local	39.2	74.2	35.9	19.8
Cohere Embed v4 alone · frontier API	53.9	77.5	40.1	20.6
Share of the gap closed by BM25	8%	39%	33%	not significant

PostgreSQL, BM25, and a 384-dimension model on one laptop, no API calls: adding keyword search closes **a third or more of the gap** to a frontier model on SciFact and NFCorpus.

Rerank: measure it too · FiQA

Candidates sent to Cohere Rerank 3.5	NDCG@10	Recall@50	p50
none: vector alone	53.9	78.5	3.5 ms
vector top 50	49.7	78.5	417 ms
RRF · BM25 top 50	50.4	75.2	1,093 ms
vector $\cup$ BM25, up to 100	49.0	76.5	1,208 ms

It lifts 403b from #26 to #2, but puts a known answer first on 47.5% of questions against 53.9% for Embed v4 alone. On no dataset did a reranked arm beat vector by more than noise.

- A reranker reads the question and each candidate **together**, but can only reorder what the lists found: Recall@50 is its ceiling.
- It runs **outside PostgreSQL**: a network call per question, and the candidates' text leaves the database.

Storage: the same vectors, smaller · FiQA

Index on <code>embedding</code>	Bytes per vector	HNSW size	NDCG@10	p50 ms
<code>vector(1536)</code>	6,148	450 MB	53.9	3.5
<code>halfvec(1536)</code> expression index	3,080	225 MB	53.7	3.4
<code>binary_quantize()</code> + rescore 200	200	28 MB	54.1	3.3

Half the index for the same NDCG within noise; 1/16 with binary plus a rescore. The table keeps full-precision vectors; only the index shrinks.

What these numbers do and don't prove

They show	They don't show
Which arm ranks judged answers higher on four public datasets	That the same holds for your data
The pitfalls, reproduced with real plans	Production latency: this is one laptop, one client
Relative cost of each arm on up to 57,638 documents	Behavior at 50 million rows or under concurrency
This embedding and rerank model pair	Future models: they improve, and results depend on the model versions used
A public benchmark	That the models never saw similar text in training
Two embedding models, one reranker	How other models behave; SCIDOCS blends are untuned (no dev split)

BEIR judgments are incomplete: an unjudged document can be a good answer and still count as a miss.

Three things to take home

**Rank words with BM25.
Blend, tuned on held-out questions.
Hybrid pays most on smaller models.**

Local model + BM25: **+1.2 to +2.1 NDCG@10 on three of four datasets.** Frontier model: equal-weight RRF never won. **Measure on your own questions before you choose.**

Take it home

The lab: `hybrid-lab/`. Numbered SQL, VS Code cells, the UI, and the scoreboard.

The skill: adds hybrid search to *your* table and measures it.

```
/plugin marketplace add shayons/talks  
/plugin install postgres-hybrid-search@shayons-talks
```

Other agents: copy `hybrid-search-
plugin/skills/postgres-hybrid-search/`.

github.com/shayons/talks · [conferences/2026-postgresconf-agentic-ai](#)



Thank you.

Which of your queries needs both words and meaning?

Shayon Sanyal · [linkedin.com/in/shayonsanyal](https://www.linkedin.com/in/shayonsanyal)

github.com/shayons/talks

Appendix · operate the measured workload

Area	Inspect before changing it
HNSW build	<code>maintenance_work_mem</code> (graph must fit), parallel workers, <code>m</code> , <code>ef_construction</code> , index size
HNSW search	<code>ef_search</code> $\geq$ LIMIT, filtered recall, <code>iterative_scan</code> , <code>max_scan_tuples</code>
Full text	OR vs AND queries, match counts, <code>ts_rank_cd</code> cost on common terms, BM25 availability
Embeddings	Model id per column, input types, re-embedding on text change, quota and throttling
Functions	Inner plans with <code>auto_explain</code> ; <code>plan_cache_mode</code> for optional filters
Evaluation	Re-run the question set after every change; keep per-question wins and losses

Appendix · FiQA scoreboard, every arm

Arm	NDCG@10 (bar) and Recall@50 (line)	Recall@50	p50 ms	vs Embed v4 vector
Embed v4 · binary + rescore	54.1	78.5	3.3	11 / 11
Tuned blend · Embed v4	53.9	78.5	8.7	0 / 0
Vector · Embed v4	53.9	78.5	3.5	baseline
Embed v4 · halfvec	53.7	78.3	3.4	2 / 3
RRF · ts_rank_cd + Rerank	51.2	72.6	1634.6	161 / 230
RRF · BM25 + Rerank	50.4	75.2	1092.5	157 / 239
Embed v4 + Rerank	49.7	78.5	416.5	158 / 241
Embed v4 ∪ BM25 + Rerank	49.0	76.5	1207.5	148 / 248
RRF · BM25 + Embed v4	41.3	75.2	8.2	98 / 325
Tuned blend · bge-small	39.2	62.3	7.3	86 / 342
Vector · bge-small (local)	38.0	59.8	2.0	81 / 344
RRF · BM25 + bge-small	34.9	61.8	6.6	73 / 381
RRF · ts_rank_cd + Embed v4	33.8	72.6	149.9	37 / 443
BM25	23.6	47.1	5.2	38 / 440
Score sum	5.6	15.8	142.3	8 / 512
Keyword · all words	4.3	6.7	0.4	6 / 506
Concatenate	2.9	14.9	150.5	6 / 518
Keyword · ts_rank_cd	2.8	14.6	165.8	6 / 518

Bar: NDCG@10 · line: Recall@50 · laptop latency, rerank rows include Bedrock calls · [results/scoreboard-fiqa.md](#)

References

- **PostgreSQL 18:** [Full-text search, Using EXPLAIN, auto_explain](#)
- **pgvector 0.8.6:** [HNSW, filtering, iterative index scans](#) · **pg_textsearch 1.4.0:** [BM25 for PostgreSQL](#)
- **RRF:** Cormack, Clarke & Büttcher, [SIGIR 2009](#) · **HNSW:** Malkov & Yashunin, [arXiv:1603.09320](#)
- **Score blending:** Bruch, Gai & Ingber, "An Analysis of Fusion Functions for Hybrid Retrieval", [ACM TOIS 2023](#) · **bge-small-en-v1.5:** [BAAI, MIT license](#), run with [fastembed](#)
- **FiQA-2018, SciFact, NFCorpus, SCIDOCS** via [BEIR](#) (Thakur et al., NeurIPS 2021) · Dave Ebbelaar, [hybrid-retrieval tutorial](#) (the FiQA + NDCG approach this lab moves into PostgreSQL)
- **This talk:** [hybrid-lab/sql/](#), [hybrid-lab/results/](#), [hybrid-search-plugin/](#)

Measured on PostgreSQL 18.6, pgvector 0.8.6, pg_textsearch 1.4.0 · September 2026