

Hybrid Search in PostgreSQL

Combining Vector and Full-Text for Real-World Applications

Shayon Sanyal, Principal Worldwide PostgreSQL Specialist SA

Postgres Summit US 2026, New York City
September 30, 10:30–11:20 EDT, Letterpress

Two questions. Two different misses.

Keyword search misses

"How do I cancel my subscription?"

The help page is titled **"How to end your membership."** Not one word in common, so keyword search never sees it. Vector search matches it on meaning.

Vector search misses

"AA batteries"

To an embedding, **AAA batteries** look almost the same: close in meaning, wrong product. Keyword search matches the exact term.

Your users type both kinds.

Today: keyword and vector search in one PostgreSQL query, and how to check on your own questions whether combining them helps.



About me

Shayon Sanyal

Principal Worldwide PostgreSQL Specialist SA

I help teams build on PostgreSQL, from relational applications to retrieval and agent workflows.

Today: readable SQL, measured results, and a skill you can point at your own tables.



[linkedin.com/in/shayonsanyal](https://www.linkedin.com/in/shayonsanyal)

How we'll measure

2,271 questions with known right answers.

FiQA-2018: 57,638 finance forum posts, 648 test questions. The main example.


SciFact, NFCorpus, SCIDOCS: chosen by one rule before measuring: public BEIR benchmark sets under 30,000 documents where keyword search (BM25) beat every dense vector model in the 2021 BEIR paper.

NDCG@10 scores a ranking higher when the known answers sit near the top.

The abstract promised that combining beats either alone. **We'll measure when it does, and when it doesn't.**

Credit for a known answer at each rank

$1 / \log_2(\text{rank} + 1)$: an answer at rank 1 is worth more than three answers at rank 10



Four datasets, four kinds of question

DATASET: WHAT IT SEARCHES	A TYPICAL TEST QUESTION	TEST QUESTIONS	KNOWN ANSWERS PER QUESTION	BM25 VS VECTOR (EMBED V4), NDCG@10
FiQA: 57,638 finance forum posts	"Where should I park my rainy-day / emergency fund?"	648	2.6	23.6 vs 53.9
SciFact: 5,183 science abstracts	"Anthrax spores can be disposed of easily after they are dispersed." A claim to check	300	1.1	68.8 vs 77.5
NFCorpus: 3,633 medical abstracts	"low-carb diets." Half are two words or fewer	323	38.2, graded	32.3 vs 40.1
SCIDOCS: 25,657 paper abstracts	"A Fast Learning Algorithm for Deep Belief Nets." A title; its answers are papers it cites	1,000	4.9	15.4 vs 20.6

Same methods, very different gaps: BM25 reaches **89%** of vector search on SciFact but **44%** on FiQA. One dataset would have told a different story.

The stack: one PostgreSQL database

PIECE	VERSION	JOB
PostgreSQL	18.6	<code>tsvector</code> , GIN, <code>ts_rank_cd</code> , the SQL that combines and scores results
pgvector	0.8.6	<code>vector(1536)</code> , HNSW, iterative index scans, <code>halfvec</code> , <code>bit</code>
pg_textsearch	1.4.0	BM25 index and <code><@></code> operator (PostgreSQL license)
Cohere Embed v4 on Amazon Bedrock	1536 dims	Posts as <code>search_document</code> , questions as <code>search_query</code>
Cohere Rerank on Amazon Bedrock	3.5	Reorders a top 50 outside the database
bge-small-en-v1.5 with fastembed	384 dims	A small open-source model, on this laptop
VS Code + SQLTools		Numbered SQL files you can run yourself

The database side is all open source: PostgreSQL, pgvector, pg_textsearch. So is bge-small, and that fully open path is where hybrid pays most. Every number today is from this laptop.

One row, three representations

01_schema.sql

```
CREATE TABLE docs (  
  id          text PRIMARY KEY,  
  body        text NOT NULL,  
  tsv         tsvector GENERATED ALWAYS AS (to_tsvector('english', body)) STORED,  
  embedding   vector(1536), -- Cohere Embed v4, input_type = 'search_document'  
  embedding_local vector(384) -- bge-small-en-v1.5, open source, local  
);  
  
CREATE INDEX ON docs USING gin (tsv); -- 0.6 s  
CREATE INDEX ON docs USING hnsw (embedding vector_cosine_ops); -- 30.4 s  
CREATE INDEX ON docs USING bm25 (body) WITH (text_config = 'english'); -- 3.9 s
```

The lexemes are **generated** and cannot drift from the text. Embeddings are not: **one column per model**, re-embedded when the text changes.

Pitfall 1: every word must match

03_keyword_and.sql

```
SELECT websearch_to_tsquery('english', 'Where should I park my rainy-day / emergency fund?');  
-- 'park' & 'rainy-day' <=> 'raini' <=> 'day' & 'emerg' & 'fund'
```

`websearch_to_tsquery` and `plainto_tsquery` join words with **AND**. A natural question rarely has all its words in one answer.

405 of 648 questions match no post at all. NDCG@10: 4.3

Match any word, rank the matches

04_keyword_or.sql

```
WITH question AS (  
  SELECT replace(plainto_tsquery('english', body)::text, ' & ', ' | ')::tsquery AS tsq  
  FROM active_query  
)  
SELECT d.id, ts_rank_cd(d.tsv, q.tsq) AS score  
FROM docs d, question q  
WHERE d.tsv @@ q.tsq  
ORDER BY score DESC  
LIMIT 50;
```

Recall@50 (known answers anywhere in the top 50) doubles, 6.7 → 14.6, but NDCG@10 **falls to 2.8**: with no IDF, “fund” and “day” outweigh “rainy-day”. The rainy-day question matches **10,460 posts**, and `ts_rank_cd` scores every one of them: **81 ms**.

ts_rank_cd is not BM25 05_bm25.sql

```
SELECT id, -(body <@> to_bm25query(:'question', 'docs_body_bm25')) AS bm25
FROM docs
ORDER BY body <@> to_bm25query(:'question', 'docs_body_bm25')
LIMIT 50;
```

BM25 weighs rare terms up (IDF), saturates repeats, and normalizes for length. The index keeps the corpus statistics and returns the top 50 **without scoring every match**. 23.6 is exactly BEIR's published BM25 score for FiQA (0.236).

KEYWORD METHOD	NDCG@10	TOP 50, RAINY-DAY QUESTION
ts_rank_cd, any word	2.8	81 ms
BM25, pg_textsearch	23.6	0.27 ms

Vector search: the top 50 by meaning `06_vector.sql`

```
SET hnsw.ef_search = 100;           -- default 40: LIMIT 50 would silently return 40
SELECT id, 1 - (embedding <=> (SELECT embedding FROM active_query)) AS cosine
FROM docs
WHERE embedding IS NOT NULL
ORDER BY embedding <=> (SELECT embedding FROM active_query)
LIMIT 50;
```

Keep the distance operator in `ORDER BY`, ascending, with a `LIMIT`. Embed questions as `search_query` and posts as `search_document`.

NDCG@10 **53.9**, median **3.5 ms** per question

Pitfall 2: adding scores from different scales

07a_naive_sum.sql

07b_concat_dedupe.sql

RAINY-DAY QUESTION, EACH LIST'S TOP 50		MIN	MAX
ts_rank_cd		1.70	5.40
cosine similarity		0.478	0.598
HOW THE TWO LISTS ARE COMBINED			NDCG@10
keyword_score + cosine			5.6
keyword list, then vector list, duplicates removed			2.9
Reciprocal Rank Fusion of the same two lists			33.8

Same inputs, 6× the score. `ts_rank_cd`'s normalization options keep its scores in a range, but still not comparable to cosine.

Combine ranks, not scores

$$\text{RRF}(d) = \sum 1 / (60 + \text{rank}_{\text{list}}(d))$$

BM25	VECTOR	COMBINED BY RRF
1 66626	1 293687	1 32811 known answer
2 371922	2 348514	2 371922 0.02912
3 569342	3 469809	3 469809 0.02778
4 32811 known answer	4 361639	4 424427 0.02633
5 2128	5 458063	5 32671 0.02568
6 32671	6 427365	6 72960 0.02452
7 404800	7 209789	7 488737 0.02379
8 266457	8 283692	8 152096 0.02114
9 272840	9 32811 known answer	9 293687 0.01639
10 65567	10 62897	10 66626 0.01639

Known answer: BM25 #4, vector #9 → $1/64 + 1/69 = \mathbf{0.03012}$, combined #1.

Each list's own #1 is missing from the other list: $1/61 = 0.01639$, combined #9 and #10.

RRF is a full outer join

08a_hybrid_rrf.sql

```
WITH keyword AS (  
    SELECT id, row_number() OVER (ORDER BY ts_rank_cd(tsv, q) DESC, id) AS rank  
    FROM docs, (SELECT replace(plainto_tsquery('english', :'question')::text,  
                                ' & ', ' | ')::tsquery AS q) question  
    WHERE tsv @@ q ORDER BY rank LIMIT 50  
,  
semantic AS (  
    SELECT id, row_number() OVER (ORDER BY distance, id) AS rank  
    FROM (SELECT id, embedding <=> :'question_embedding' AS distance  
          FROM docs WHERE embedding IS NOT NULL  
          ORDER BY distance LIMIT 50) nearest  
)  
SELECT coalesce(k.id, v.id) AS id,  
       coalesce(1.0 / (60 + k.rank), 0) + coalesce(1.0 / (60 + v.rank), 0) AS rrf  
FROM keyword k FULL OUTER JOIN semantic v USING (id)  
ORDER BY rrf DESC  
LIMIT 10;
```

One statement, two indexes (GIN and HNSW). Median **150 ms** with `ts_rank_cd`; **8.2 ms** with BM25 (`08b_hybrid_rrf_bm25.sql`). The two lists run **one after the other** in one backend process, not in parallel: the plan has no Gather node, and 8.2 ms is about BM25's 5.2 plus vector's 3.5.

Blend scores the right way

08c_hybrid_blend.sql

```
keyword_norm AS (
    -- per question: min-max to [0, 1]; semantic_norm is the same
    SELECT id, (score - min(score) OVER ()) / (max(score) OVER () - min(score) OVER ()) AS s
    FROM keyword
)
SELECT id, w * coalesce(v.s, 0) + (1 - w) * coalesce(k.s, 0) AS blend
FROM keyword_norm k FULL OUTER JOIN semantic_norm v USING (id);
```

RRF ignores how **strong** each match is. A blend keeps it: scale each list's scores to 0–1, then weight them. It is what “adding scores” should have been in pitfall 2.

WEIGHT w ON VECTOR, CHOSEN ON SEPARATE TUNING QUESTIONS	FIQA	SCIFACT	NFCORPUS	SCIDOCs
Cohere Embed v4 (frontier)	1.0	0.70	0.70	0.5, not tuned
bge-small (local)	0.65	0.55	0.65	0.5, not tuned

Knobs that change different things

KNOB	CONTROLS	WATCH
Results per list (50)	How many posts each list contributes	A post outside both lists can't be combined or reranked: Recall@50
k (60)	How fast credit falls with rank	Measured k = 5 to 200: at most 4.3 NDCG@10 (FiQA), 1.4 elsewhere. No k made equal-weight RRF beat Embed v4 alone
Weights	Trust in each list	Equal-weight RRF lost to Embed v4 alone on all four datasets; tune a blend's weight on separate questions instead
<code>hnsw.ef_search</code> (100)	How hard each HNSW scan searches	Must be $\geq$ results per list

Change one at a time, and re-run the test questions.

Live on NFCorpus: all local except the last column

Small model + BM25 beat the frontier model here

NFCorpus test question PLAIN-307, with 3 known answers, also active in VS Code

Vitamin D: Shedding some light on the new recommendations

A in 4 of 4 top 10s **B** in 3 of 4 top 10s **C** in 4 of 4 top 10s

BM25

SQL

65

NDCG@10

3 of 3 answers in the top 10
10.5 ms, live SQL

- 1 Safety considerations and potential interactions of vitamins: should vitamins be considered drugs?...
- 2 **B** VITAMIN D: A D-LIGHTFUL SOLUTION FOR HEALTH Throughout evolution sunlight produced vitamin D in the ski...
- 3 The aryl hydrocarbon receptor and its xenobiotic ligands: a fundamental

Vector (bge-small, local)

SQL

77

NDCG@10

2 of 3 answers in the top 10
12.8 ms, live SQL

- 1 **A** Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 2 **C** Evaluation, treatment, and prevention of vitamin D deficiency: an Endocrine Society clinical practice guideline....
- 3 Traditionally living populations in East

Tuned blend (bge-small)

SQL

97

NDCG@10

3 of 3 answers in the top 10
15.7 ms, live SQL

- 1 **C** keyword 4 vector 2 Evaluation, treatment, and prevention of vitamin D deficiency: an Endocrine Society clinical practice guideline....
- 2 **A** keyword 8 vector 1 Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 3 keyword 14 vector 4

Vector (Embed v4)

SQL

71

NDCG@10

3 of 3 answers in the top 10
21.5 ms, live SQL

- 1 Low Vitamin D Status: Definition, Prevalence, Consequences and Correction Vitamin D is obtained fro...
- 2 **A** Towards prevention of vitamin D deficiency and beyond: knowledge gaps and research needs in vitamin D...
- 3 **B** VITAMIN D: A D-LIGHTFUL SOLUTION

localhost:8018, pick NFCorpus, then “Local hybrid beats the frontier model”

From text to vector

- **Cohere Embed v4** on Amazon Bedrock: 1536 dimensions, unit-length vectors, so cosine distance and inner product rank the same way.
- **Two input types.** Posts are embedded as `search_document`, questions as `search_query`. Swapping them lowers relevance without any error.
- **One model per column.** Vectors from different models or versions are not comparable, even at the same dimension. Store the model id.
- **A small local model too:** bge-small-en-v1.5, 384 dims, open source, in its own column. FiQA's 57,638 posts took about an hour on the laptop CPU; Embed v4 took 35 minutes behind a tokens-per-minute quota. Either way, a rerun fills only the rows still NULL.

Pitfall 3: a filter is not a pre-filter

09_filtered_hybrid.sql

```
SELECT id FROM docs
WHERE tsv @@ plainto_tsquery('english', 'mortgage') -- 4.2% of posts
ORDER BY embedding <=> : 'question_embedding' LIMIT 10;
```

SETTING	PLAN	ROWS RETURNED
defaults (<code>iterative_scan = off</code> , <code>ef_search</code> <code>= 40</code>)	HNSW, then filter: Rows Removed by Filter: 39	1 of 10
<code>SET hnsw.iterative_scan = relaxed_order</code>	HNSW keeps walking until 10 pass	10 of 10
a rare term: <code>heloc</code> (0.2%)	GIN, then exact sort	10 of 10

pgvector 0.8 iterative scans are bounded by `hnsw.max_scan_tuples` (20,000).

Similar to the question AND contains a word

11_hybrid_function.sql

```
SELECT h.*, left(d.body, 60)
FROM hybrid_search(:'question', :'question_embedding',
                  match_count => 5, required_terms => 'mortgage') h
JOIN docs d ON d.id = h.doc_id;
```

```
CREATE FUNCTION hybrid_search(...) ... LANGUAGE sql STABLE
SET hnsf.ef_search = 100
SET hnsf.iterative_scan = relaxed_order
SET plan_cache_mode = force_custom_plan -- keeps "required_terms IS NULL OR ..." indexable
```

The required word goes into **both** lists, before their LIMIT.

Boost and expand: two patterns on top of `hybrid_search()`

Boost by recency, popularity, preference

```
SELECT h.doc_id,
       h.score
  * power(0.5, extract(epoch FROM
                      now() - d.published_at) / 86400 / 30)
  * (1 + ln(1 + d.likes) / 10)
  * CASE WHEN d.category = ANY(:'preferred')
        THEN 1.2 ELSE 1 END AS boosted
FROM hybrid_search(:'question',
                  :'question_embedding', match_count => 50) h
JOIN docs d ON d.id = h.doc_id
ORDER BY boosted DESC
LIMIT 10;
```

Expand along links, two hops

```
WITH RECURSIVE related (doc_id, score, depth) AS (
  SELECT doc_id, score, 0
    FROM hybrid_search(:'question',
                      :'question_embedding', match_count => 10)
  UNION ALL
  SELECT l.dst, r.score / 2, r.depth + 1
    FROM related r
   JOIN links l ON l.src = r.doc_id
  WHERE r.depth < 2
)
SELECT doc_id, max(score) AS score
  FROM related
 GROUP BY doc_id
 ORDER BY score DESC
LIMIT 10;
```

Boost a larger pool than you show (50, then keep 10), or a boost can't lift anything. **Not measured here:** these datasets have no dates, popularity, or links.

Hybrid vs vector, same embedding model: NDCG@10 change on 2,271 test questions

	FIQA	SCIFACT	NFCORPUS	SCIDOCS
Small local model: bge-small, 384 dims, on this laptop				
Tuned blend + BM25	+1.2 ↑ +0.4...+2.1	+2.1 ↑ +0.0...+4.3	+2.1 ↑ +1.3...+3.0	+0.2 * -0.5...+0.8
Equal-weight RRF + BM25	-3.1 ↓ -4.7...-1.5	+1.7 -0.9...+4.3	+2.3 ↑ +1.0...+3.7	-0.3 -1.0...+0.5
Frontier model cut to its first 256 dims: Cohere Embed v4				
Tuned blend + BM25	+0.5 -0.0...+1.1	+3.5 ↑ +1.3...+5.7	+2.8 ↑ +1.5...+4.1	+1.9 ↑ * +1.3...+2.5
Frontier model: Cohere Embed v4, 1536 dims				
Tuned blend + BM25	0.0 chose vector only	+0.2 -1.2...+1.4	+0.8 ↑ +0.1...+1.5	-0.8 ↓ * -1.4...-0.2
Equal-weight RRF + BM25	-12.5 ↓ -14.5...-10.6	-2.9 ↓ -5.1...-0.6	-1.1 -2.2...+0.1	-1.2 ↓ -2.0...-0.5
Cohere Rerank 3.5, top 50	-4.1 ↓ -6.0...-2.3	-0.4 -2.8...+1.9	-1.9 ↓ -3.3...-0.4	-0.4 -1.1...+0.3

↑ / ↓: 95% confidence range entirely above / below zero, so not noise. * SCIDOCS: no tuning questions, weight 0.5

When hybrid pays: a smaller model

NDCG@10	FIQA	SCIFACT	NFCORPUS	SCIDOCs
bge-small alone (local, open source)	38.0	72.0	33.8	19.6
bge-small + BM25, tuned blend (all local)	39.2	74.2	35.9	19.8
bge-small + <code>ts_rank_cd</code> , tuned blend (core PostgreSQL only)	37.9	72.4	34.4	16.0 ↓
Cohere Embed v4 alone (frontier API)	53.9	77.5	40.1	20.6
Share of the gap closed by BM25	8%	39%	33%	within noise

PostgreSQL, BM25, and a 384-dimension model on one laptop, no API calls: adding keyword search closes **a third or more of the gap** to a frontier model on SciFact and NFCorpus. Without BM25's IDF, `ts_rank_cd` gained nothing beyond noise.

Rerank on FiQA: measure it too

WHAT COHERE RERANK 3.5 REORDERS	NDCG@10	RECALL@50	MEDIAN
none: vector alone	53.9	78.5	3.5 ms
vector top 50	49.7	78.5	417 ms
RRF (BM25) top 50	50.4	75.2	1,093 ms
vector $\cup$ BM25, up to 100	49.0	76.5	1,208 ms

It lifts 403b from #26 to #2, but puts a known answer first on 47.5% of questions against 53.9% for Embed v4 alone. On no dataset did a reranked method beat vector search by more than noise.

- A reranker reads the question and each result **together**, but can only reorder what the lists found: Recall@50 is its ceiling.
- It runs **outside PostgreSQL**: a network call per question, and the results' text leaves the database.

Storage on FiQA: fewer bits or fewer dimensions?

INDEX ON EMBEDDING	BYTES PER VECTOR	HNSW SIZE	NDCG@10	MEDIAN MS
vector(1536)	6,148	450 MB	53.9	3.5
halfvec(1536) expression index	3,080	225 MB	53.7	3.4
binary_quantize(), then re-sort 200 by full vector	200	28 MB	54.1	3.3
first 1024 dims, subvector()	4,104	450 MB	53.4	2.9
first 512 dims	2,056	150 MB	51.8 ↓	2.2
first 256 dims	1,032	75 MB	49.3 ↓	1.7

Quantize before you truncate. Fewer bits kept NDCG within noise; 512 and 256 dims lost 1.6 to 5.5 points on every dataset. And 1024 dims saves no index space: an 8 KB page still holds one vector. If you do truncate, add BM25: at 256 dims it won back 46% to 64% of the loss on three datasets.

What these numbers do and don't prove

THEY SHOW	THEY DON'T SHOW
Which method ranks known answers higher on four public datasets	That the same holds for your data
The pitfalls, reproduced with real plans	Production latency: this is one laptop, one client
Relative cost of each method on up to 57,638 documents	Behavior at 50 million rows or under concurrency
This embedding and rerank model pair	Future models: they improve, and results depend on the model versions used
A public benchmark	That the models never saw similar text in training
Two embedding models, one reranker	How other models behave; SCIDOCS blends aren't tuned (no tuning questions)
English questions, word matching	Multilingual queries or phrase search (<->)
Boost and link-expansion patterns	Their effect: these datasets have no dates, popularity, or links

The judgments are incomplete: an unjudged document can be a good answer and still count as a miss.

Three things to take home

Rank words with BM25.
Blend, tuned on separate questions.
Hybrid pays most on smaller models.

Local model + BM25: **+1.2 to +2.1 NDCG@10 on three of four datasets.** Frontier model: equal-weight RRF never won, not even on FiQA's 177 questions with a number or acronym (vector 57.7, RRF 46.5). **Measure on your own questions before you choose.**

Take it home

The lab: `hybrid-lab/`. Numbered SQL, VS Code cells, the UI, and the scoreboard.

The skill: adds hybrid search to *your* table and measures it.

```
/plugin marketplace add shayons/talks  
/plugin install postgres-hybrid-search@shayons-talks
```

Other agents: copy `hybrid-search-
plugin/skills/postgres-hybrid-search/`.

github.com/shayons/talks, in [conferences/2026-postgresconf-agentic-ai](#)



Thank you.

Which of your queries needs both words and meaning?



Shayon Sanyal

[linkedin.com/in/shayonsanyal](https://www.linkedin.com/in/shayonsanyal)

github.com/shayons/talks

References

- **PostgreSQL 18:** [Full-text search](#), [Using EXPLAIN](#), [auto_explain](#)
- **pgvector 0.8.6:** [HNSW](#), [filtering](#), [iterative index scans](#); **pg_textsearch 1.4.0:** [BM25 for PostgreSQL](#)
- **RRF:** Cormack, Clarke & Büttcher, [SIGIR 2009](#); **HNSW:** Malkov & Yashunin, [arXiv:1603.09320](#)
- **Score blending:** Bruch, Gai & Ingber, "An Analysis of Fusion Functions for Hybrid Retrieval", [ACM TOIS 2023](#); **bge-small-en-v1.5:** [BAAI](#), [MIT license](#), run with [fastembed](#)
- **FiQA-2018, SciFact, NFCorpus, SCIDOCS** via [BEIR](#) (Thakur et al., NeurIPS 2021); Dave Ebbelaar, [hybrid-retrieval tutorial](#) (the FiQA + NDCG approach this lab moves into PostgreSQL)
- **This talk:** [hybrid-lab/sql/](#), [hybrid-lab/results/](#), [hybrid-search-plugin/](#)

Measured on PostgreSQL 18.6, pgvector 0.8.6, pg_textsearch 1.4.0, September 2026